

Language support for AOP

AspectJ and beyond

Mario Südholt

www.emn.fr/sudholt

INRIA and École des Mines de Nantes

OBASCO project, Nantes, France

Plan

- AOP and AspectJ
- AspectJ: end of story?
- Beyond AspectJ

I. AOP and AspectJ

Defining characteristics of AOP?

- **Quantification:** modularization of crosscutting concerns
- **Obliviousness:** non-anticipation; incremental development

⇒ Tackle crosscutting in large-scale applications throughout the software life cycle

More probably later from Bob . . .

What's new? (1)

- What about **computational reflection**?
 - 3-Lisp, CLOS, Reflex [Tanter et al., OOPSLA'03], ...
 - General enough reflective system can “emulate” AOP systems
 - Difficult to understand
 - Performance issues
 - Semantics issues, lack of correctness guarantees

What's new? (2)

- What about **transformation systems**?
 - General enough transformation system can “emulate” AOP systems
 - SOOT, Recoder, CIL, ...
 - Difficult to understand
 - Correctness properties difficult to handle

Yes, it is! (in a sense)

Goals for AOP

- Provide abstractions general enough to modularize (some or all) concerns.
- Be specific enough to make such modularization understandable, tractable and amenable to testing, analysis, verification of properties.

AspectJ in one slide

Base program: `critical, access`

AspectJ in one slide

```
pointcut critAcc(Base r): call(void Base.acc(int)
    && target(r)
    && cflow(call(void Base.crit(int))));
```

AspectJ in one slide

```
pointcut critAcc(Base r): call(void Base.acc(int)
    && target(r)
    && cflow(call(void Base.crit(int))));
void around(Base r): critAcc(r) {
    calls++;
    if (ok()) proceed(r);
}
```

AspectJ in one slide

```
aspect ProfBar pertarget call(void Base.acc(int)) {  
    int calls = 0;  
  
    pointcut critAcc(Base r): call(void Base.acc(int)  
        && target(r)  
        && cflow(call(void Base.crit(int))));  
  
    void around(Base r): critAcc(r) {  
        calls++;  
        if (ok()) proceed(r);  
    }  
}
```

AspectJ in one slide

```
aspect ProfBar pertarget call(void Base.acc(int)) {  
    int calls = 0;  
    static int Base.calls = 0;  
  
    pointcut critAcc(Base r): call(void Base.acc(int)  
        && target(r)  
        && cflow(call(void Base.crit(int))));  
  
    void around(Base r): critAcc(r) {  
        calls++;  
        if (ok()) proceed(r);  
    }  
}
```

Characteristics of AspectJ

- + Join points
- + Pointcuts
- + Advice
- + Aspects
- + Inter-type declarations
- +− Aspect instantiation (coarse-grained)
- +− Aspect activation (on/off)
- +− Aspect composition (dominate)

II. AspectJ: end of story?

- Other characteristics of aspect languages
- Other base languages, execution environments
- More expressive pointcut languages

Other characteristics of aspect languages

- Aspect instantiation
E.g., runtime instances, Kevin's talk
- Aspect activation
E.g., enable/disable aspects at runtime
- Aspects of aspects
E.g., layered aspects, Kevin's talk
- Aspect composition
E.g. for conflict resolution
- Weaver semantics
E.g., no aspects of aspects

A world outside Java?

- Crosscutting concerns in large (legacy) C applications
- Ex.: optimization of web caches without cache flushes
- New aspect languages for expression of complex context conditions

Other pointcut languages (1)

- **Stateful pointcuts** (explicit state in pointcuts)
 - Sequence pointcuts:
Ex.: protocol translation and bug correction
 - Temporal logic pointcuts:
Ex.: manipulation of Linux kernel code
 - Regular expression pointcuts:
Enable interference analysis among aspects

Other pointcut languages (2)

- AOP and distributed applications
 - Often integration/configuration of existing distribution platforms (see Kevin's talk)
⇒ distribution implicit to aspects
 - **Remote pointcuts** [Nishizawa et al., AOSD'04]: explicit hosts, advice server
 - Trade-off: hide complexity vs. flexibility
- **Data-flow pointcuts** [Masuhara, Kawauchi; APLAS'03], e.g., for security enforcement.
Efficiency realization

III. Beyond AspectJ

1. Dynamic aspects for C system-level applications
2. Temporal logic pointcuts for Linux kernel evolution

1. Dynamic aspects for C system-level applications

- Software evolution frequently to be performed on running systems (e.g., high-availability servers)
- Ex. concerns in a web cache
 - Modification of caching policies
 - Optimizations (e.g., protocol transformations TCP→UDP)
 - Bug corrections
- Some large applications:
Open-source web-cache “squid”: 9 MB of source

Ex.: explicit sequences for buffer overflows

- Aspect language with explicit sequences

```
seq( call(void* malloc(size_t))  
      && args(allocatedSize) && return(buffer) ;  
  write(buffer) && size(writtenSize)  
    && if(writtenSize > allocatedSize)  
      then reportOverflow(); *  
  call(void free(void*)) )
```

Aspect language

- Primitive pointcuts: calls and variables accesses (to global and local variables)
- `cflow` for nested calls (like AspectJ)
- Sequences with
 - Conditionals over data
Principally equalities (e.g. over file handles)
 - Means for ressource handling
Optimize ressource usage (e.g., reuse of file handles)

Realization: the Arachne system

- Dynamic aspect application for C without program interruption
www.emn.fr/x-info/arachne
- Rewrite binary code on the fly to weave (and deweave) aspects
- Current weaving semantics excludes nested aspects
Simplified implementation, somewhat more efficient
- [Ségura et al, AOSD'03] [Fritz et al, AOSD'05]

2. Temporal logic pointcuts for Linux kernel evolution


- Problem: support extensions of the Linux kernel over a range of kernel versions
E.g., over one major version number
- Ex.: support application-specific schedulers
E.g., for multi-media streaming
- Context: integrate an existing system for scheduler development with the kernel

Bossa: new schedulers for plain old Linux

- Bossa: system for scheduler development
www.emn.fr/x-info/bossa
- DSL: definition of scheduling policies
- Support runtime for hierarchical, prioritized, etc., schedulers
- Runtime overhead $< 5\%$

Bossa architecture

- Events mediate between (instrumented) kernel and Bossa runtime, which supports policies

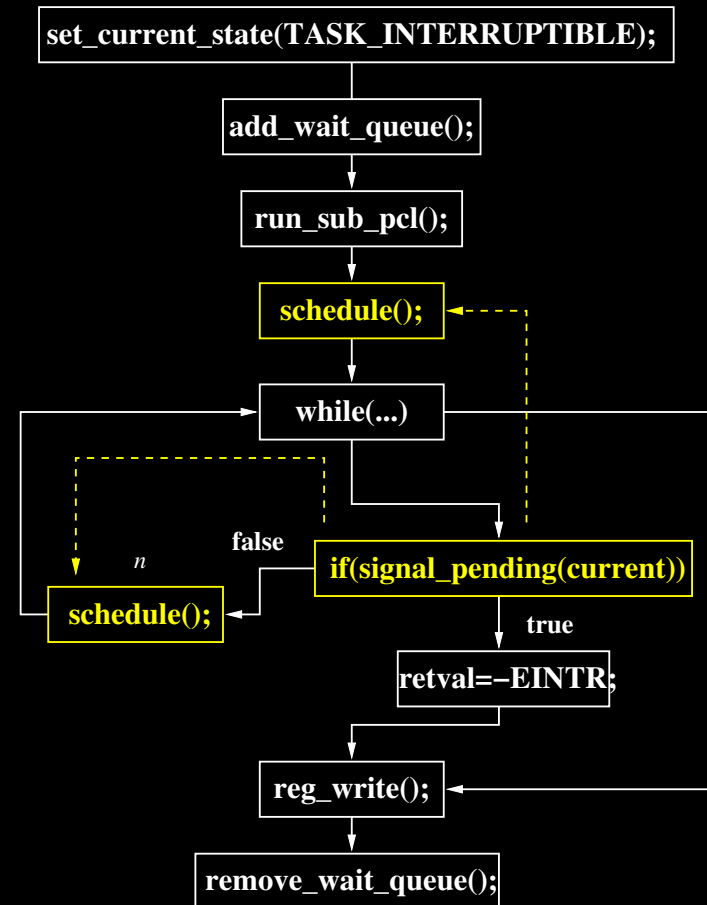


Mediation through events crosscuts the kernel

- Instrument kernel code + drivers (~ 100 MB source code)
- Instrumentation for Bossa:
 - ~ 400 instructions changed in about 150 files
- Previously manually done for Linux kernel 2.4
- Can we do better with aspects?

Problem: context dependencies

- Generate events for schedule instructions
- Other instructions relevant (e.g., thread state, yield)
- Problem: thread context implicit
- Explicit context dependencies vs. efficiency



Solution: temporal logic pointcuts

- Use temporal predicates to express control-flow relationships

n : Rewrite(n , schedule_running)

If $n \vdash AX \triangle (A \triangle (\neg \text{changeOfState}() \text{ U } \text{changeToRunning}()))$

“Change current instruction to schedule_running if for all backward pathes starting from the predecessor node, all backward pathes change to running without previous changes to the state.”

Results

- Transformational system for Bossa integration: 25 rules
- Implementation based on CIL yields exact instrumentation
⇒ no overhead to manual instrumentation
- 6 bugs of manual instrumentation detected
- [Åberg et al., ASE'03]

Conclusion

- AOP is relevant to software development
- AOP interesting from theoretical and practical viewpoint
- AspectJ is an interesting language and tool but not the end of the story

Future work

- (Almost) everything still to be done
- AOP for distributed programming
 - Remote pointcut: extend language, implementation, remote aspect calculus
- Aspect interactions
 - Generalize first results over regular expressions, use of model checking
- Aspects and components
 - Aspects over components with explicit protocols